



Universidade Estadual de Campinas

School of Electrical and Computer
Engineering (FEEC)



Seminars in Computer Engineering

IA382A - 1S2026

Automated Knowledge Artifact Creation: Using Multiple AI Tools to Explore Privacy and Data Protection in AI Systems

Final Class Project

Name: Felipe Rezende Gardin

RA: 215766

June 2026

Contents

1	Introduction	4
1.1	Background: Neeli Prasad’s Seminar on Human-AI Col- laboration	4
1.2	Why Privacy in AI Systems	5
1.3	Project Objective	5
1.4	The Four-Tool Pipeline	5
1.5	Document Structure	6
2	Methodology	6
2.1	Phase 1: Perplexity (Research)	6
2.1.1	Objective	6
2.1.2	Search Queries	6
2.1.3	Configuration	7
2.1.4	Output Summary	7
2.2	Phase 2: Claude API (Synthesis)	8
2.2.1	Objective	8
2.2.2	Model and Parameters	9
2.2.3	System Prompt	9
2.2.4	User Prompt	9
2.2.5	Execution Notes	9
2.3	Phase 3: QuillBot (Refinement)	10
2.3.1	Objective	10
2.3.2	Tool Configuration	10
2.3.3	Input Text Sample	10
2.3.4	Output Variations	11
2.3.5	Limitations Encountered	11
2.4	Phase 4: SlidesAI.io (Presentation)	12
2.4.1	Objective	12
2.4.2	Tool Configuration	12

2.4.3	Input Preparation	12
2.4.4	Output Description	12
2.5	Execution Timeline and Resource Usage	13
3	Results	14
3.1	Research Findings (Perplexity Output)	14
3.2	Synthesized Content (Claude API Output)	14
3.2.1	Section A: Why Privacy Matters in AI	14
3.2.2	Section B: Real Harms and the Cost of Privacy Failure	15
3.2.3	Section C: Regulatory Frameworks	15
3.2.4	Section D: Technical Safeguards	15
3.2.5	Section E: Governance and the Role of Human Oversight	15
3.3	Writing Style Comparison (QuillBot Output)	16
3.3.1	Original (Claude API)	16
3.3.2	Standard Mode (QuillBot)	16
3.3.3	Formal Mode (QuillBot)	16
3.3.4	Analysis	17
3.4	Presentation Output (SlidesAI.io)	17
3.5	Comparative Analysis: Tool Quality and Efficiency . . .	18
3.5.1	Quality Assessment	18
3.5.2	Efficiency	19
3.5.3	Observed Bottlenecks	19
3.6	Pipeline Effectiveness	19
4	Conclusions	19
4.1	Key Accomplishments	19
4.2	Limitations and Trade-Offs	20
4.3	Broader Lessons on Human-AI Collaboration	20
4.4	Implications for Privacy in AI	20

4.5	Future Work	21
4.6	Final Reflection	21
5	References	22

Abstract

This project ran four AI tools in sequence to create a knowledge artifact on Privacy and Data Protection in AI Systems, a core topic from Neeli Prasad’s IEEE Distinguished Lecture on Human-AI Collaboration. The workflow was: Perplexity gathered research on privacy risks and regulatory frameworks. Claude API synthesized the findings into a 1800-word article. QuillBot rewrote the content in different styles to compare clarity. SlidesAI.io generated a 9-slide presentation from the curated material. Every prompt, configuration, and output is documented so the process can be reproduced and examined.

Claude API produced the clearest technical writing. QuillBot highlighted how changing vocabulary affected readability. SlidesAI.io created slides with no manual design work needed. The free-tier versions of these tools did impose limits: QuillBot’s 125-word paraphrase cap and SlidesAI.io’s template rigidity. This report documents the full methodology, compares tool quality across their different roles, and reflects on what it means to use multiple AI systems responsibly in academic work. The approach mirrors Prasad’s call for transparency: showing all the steps, not just the final product.

1 Introduction

1.1 Background: Neeli Prasad’s Seminar on Human-AI Collaboration

Neeli Prasad’s IEEE Distinguished Lecture walked through three phases of AI: perception AI (medical imaging, autonomous vision systems), generative AI (large language models, code generation), and physical AI (humanoid robots). She said ethics, privacy, and governance must be built in from the start [3]. Privacy especially. It’s not something you add later.

She gave examples of what happens when privacy fails. A Hong Kong bank lost millions in a deepfake fraud where AI impersonated a CEO’s voice and video. Patient records have been leaked and used for blackmail. Images and audio of minors have been faked to harass them, linked to documented suicides. Without privacy protection, AI causes real damage.

1.2 Why Privacy in AI Systems

The topic feels urgent right now. The EU updated the AI Act in 2024 [2]. GDPR was 2016 [1]. That ten-year lag suggests AI privacy law is moving faster now. Anyone entering engineering should know this. It's not optional.

Prasad also pushed for bringing different disciplines together. Privacy isn't just a technical problem. You need engineers, lawyers, ethicists, psychologists in the room at once, thinking about consent, data ownership, surveillance, labor displacement, and human dignity.

1.3 Project Objective

This project takes a practical approach. Rather than write an essay, it runs privacy research through four AI tools in sequence. Perplexity gathers research. Claude synthesizes it. QuillBot refines the writing. SlidesAI.io generates slides. The whole process is documented, including every prompt and output [3]. This follows Prasad's insistence that students show their prompts and reasoning.

1.4 The Four-Tool Pipeline

The pipeline works in four stages:

1. **Perplexity** (Research): Gather cited information on AI privacy risks, regulatory frameworks, and real-world cases.
2. **Claude API** (Synthesis): Analyze the seminar content and research findings, then generate an article on privacy mechanisms, compliance challenges, and technical safeguards.
3. **QuillBot** (Refinement): Rewrite content in different styles to see how tone and wording affect clarity.
4. **SlidesAI.io** (Presentation): Turn the curated content into a slide deck.




Figure 1: The four-stage pipeline: Perplexity (research) → Claude API (synthesis) → QuillBot (refinement) → SlidesAI.io (presentation)

1.5 Document Structure

Section 2 covers the detailed prompts, settings, and execution for each tool. Section 3 shows the actual outputs and compares their quality and efficiency. Section 4 concludes with what was learned and connects back to Prasad’s themes: transparency and governance in human-AI collaboration.

2 Methodology

This section documents every step of the four-tool pipeline: the exact prompts, configurations, and execution parameters used at each stage. The goal is reproducibility and transparency. Following Prasad’s call to “disclose your prompt engineering,” we show the work, not just the results.

2.1 Phase 1: Perplexity (Research)

2.1.1 Objective

Gather structured, peer-reviewed and cited information on AI privacy risks, regulatory frameworks (GDPR, EU AI Act), real-world breach cases, and technical safeguards such as differential privacy and privacy-by-design principles.

2.1.2 Search Queries

The following queries were executed sequentially on Perplexity (<https://www.perplexity.ai/>):

1. “What are the main privacy risks of AI systems in 2025-2026?”
2. “What is Privacy by Design and how does it apply to AI systems?”

3. ‘‘Key provisions of the EU AI Act related to data protection and privacy’’
4. ‘‘Real-world cases of AI privacy violations and data breaches’’
5. ‘‘Differential privacy and federated learning as technical safeguards’’

2.1.3 Configuration

- Focus Mode: Academic (prioritizes peer-reviewed sources and research)
- Language: English
- Output Format: Structured summaries with inline source references
- Search Depth: Standard (responses returned as organized thematic outlines)

2.1.4 Output Summary

Perplexity produced five responses. The privacy risks response covered collection and re-identification, surveillance and function creep, sensitive attribute inference, model inversion attacks, weak data sharing practices, and consent gaps. Privacy by Design was explained with concrete steps for data minimization, secure model design, transparency, and vendor controls. On the EU AI Act, it discussed risk classification, documentation requirements for high-risk systems, human oversight, and how the Act aligns with GDPR.

Real-world cases included Clearview AI scraping facial images without consent, OpenAI’s 2023 ChatGPT breach that exposed billing data and chat histories, facial recognition deployments used for surveillance, and widespread child data collection on platforms. Technical safeguards covered differential privacy, which adds gradient noise to prevent membership inference and quantifies privacy budgets, and federated learning, which uses local training with aggregated-update sharing. The two are often used together to create layered protection.




Figure 2: Perplexity search interface showing query and cited results for AI privacy risks

2.2 Phase 2: Claude API (Synthesis)

2.2.1 Objective

Take the research findings from Perplexity and synthesize them into a comprehensive, cohesive article on Privacy and Data Protection in AI Systems, incorporating themes from Prasad's lecture.

2.2.2 Model and Parameters

- Model: claude-sonnet-4-6
- Effort: High
- Thinking Enabled: false

2.2.3 System Prompt

“You are an expert in artificial intelligence ethics, privacy law, and data protection. Your task is to synthesize research on AI privacy into a structured, academic article suitable for a technical report. Focus on: (1) the nature and severity of privacy risks in modern AI systems; (2) regulatory frameworks (GDPR, EU AI Act) and their implications; (3) technical safeguards (differential privacy, federated learning, privacy-by-design); (4) real-world cases of privacy violations and lessons learned; (5) the role of transparency and human oversight in protecting privacy. Write for an audience of computer engineering students. Cite sources where possible. Emphasize that privacy is not optional but foundational to ethical AI deployment.”

2.2.4 User Prompt

“Using the research attached (from Perplexity), write a 1500-2000 word article on Privacy and Data Protection in AI Systems. Structure it as: (A) Introduction—why privacy is critical now; (B) Privacy Risks—real cases and technical vulnerabilities; (C) Regulatory Landscape—GDPR and EU AI Act; (D) Technical Safeguards—privacy-by-design, differential privacy, federated learning; (E) Governance and Human Oversight, interdisciplinary collaboration and transparency. Emphasize practical implications for AI engineers and connect to the theme of human-AI collaboration.”

[Perplexity research results inserted here—approximately 3000 words of research text]

2.2.5 Execution Notes

Claude generated the response in 45 seconds via the Claude.ai web interface. The output was 1800+ words of structured analysis on privacy in AI systems. This synthesis fed into both the refinement and presentation stages.



Figure 3: Claude API response to the synthesis prompt, showing structured article output

2.3 Phase 3: QuillBot (Refinement)

2.3.1 Objective

Explore how alternative writing styles affect clarity, accessibility, and audience engagement. QuillBot’s paraphrasing modes enable rewriting the same content in different tones without changing the underlying meaning.

2.3.2 Tool Configuration

- Tool: QuillBot Paraphraser (<https://quillbot.com/>)
- Modes Tested: Standard, Formal, Creative
- Focus: Privacy and Data Protection (selected 3-4 paragraphs from Claude output)
- Creativity Level: Balanced (moderate rephrasing)

2.3.3 Input Text Sample

For comparison, the following paragraph from the Claude API output was processed through QuillBot:

“Privacy is a fundamental right that must be protected in AI systems. As AI becomes increasingly embedded in sensitive domains such as healthcare, finance, and criminal justice, the risks of misuse grow exponentially. Data minimization, purpose limitation, and consent mechanisms are core principles that must be enforced through both technical design and regulatory oversight. Without these protections, individuals face risks ranging from discrimination and surveillance to identity theft and blackmail.”

2.3.4 Output Variations

QuillBot produced three alternative versions:

1. **Standard Mode:** Rewrites with neutral, accessible language, good for general audiences.
2. **Formal Mode:** Emphasizes academic vocabulary and complex sentence structures, suited for peer review.
3. **Creative Mode:** Introduces variation and stylistic elements, less suitable for technical reporting but useful for outreach writing.

2.3.5 Limitations Encountered

The free tier of QuillBot imposes a 125-word limit per paraphrase, necessitating multiple passes for longer text. Additionally, the tool does not preserve citations or complex formatting, requiring manual reintegration of sources.




Figure 4: QuillBot paraphraser in Standard mode showing the rewritten paragraph

2.4 Phase 4: SlidesAI.io (Presentation)

2.4.1 Objective

Automatically generate a presentation deck from the curated and refined content, enabling rapid dissemination to diverse audiences.

2.4.2 Tool Configuration

- Tool: SlidesAI.io (<https://www.slidesai.io/>)
- Input Format: Plain text (Claude API output)
- Template: Professional (clean, minimal design)
- Number of Slides: Auto-generated (approximately 8-10 slides)
- Language: English
- Slide Ratio: 16:9 (standard widescreen)

2.4.3 Input Preparation

The Claude API synthesis was formatted as a structured outline with section headings and bullet points:

- Privacy in AI Systems: What's at Stake?
- Real-World Privacy Violations in AI
- Regulatory Frameworks: GDPR and the EU AI Act
- Technical Safeguards: Privacy by Design
- Governance and Transparency in Human-AI Collaboration

2.4.4 Output Description

SlidesAI.io processed the input and generated a 9-slide presentation with:

- Title slide (with topic and author)
- Overview slide (key statistics and motivation)

- Five content slides (one per section)
- Conclusion slide (summary and key takeaways)

SlidesAI.io selected images automatically and scaled text for readability across all slides. No manual design work was needed. The tool’s customization options were available but not necessary for our workflow.




Figure 5: SlidesAI.io generated slide deck overview (9 slides)

2.5 Execution Timeline and Resource Usage

Tool	Time (minutes)	Cost	Limiting Factors
Perplexity	10–15	Free	Citation count, search depth
Claude API	3–5	Free (with Claude.ai)	Response length, context window
QuillBot	8–10	Free (limited)	125-word paraphrase limit
SlidesAI.io	10–15	Free (basic)	Template flexibility, design control
Total	31–45	Free	

Table 1: Execution time and resource usage for the complete pipeline.

3 Results

3.1 Research Findings (Perplexity Output)

Perplexity returned structured research on privacy in AI spanning 15 verified sources. Key findings included:

- **Privacy Risks:** Data minimization failure, purpose limitation violations, consent bypassing, and unauthorized secondary use of training data are the primary privacy risks in modern AI systems.
- **Real-World Cases:** The Hong Kong deepfake fraud (2023) and documented patient data blackmail show that AI privacy breaches have immediate financial and social consequences.
- **Regulatory Status:** The EU AI Act (effective 2024) mandates privacy impact assessments for high-risk AI systems. GDPR (since 2018) establishes baseline data protection but was not designed for modern AI training and fine-tuning practices.
- **Technical Mechanisms:** Differential privacy (Dwork et al., 2014) and federated learning represent mature approaches to privacy-preserving AI, though adoption remains limited in production systems.
- **Governance Gap:** Technical privacy mechanisms exist, but governance structures and enforcement mechanisms lag behind. This gap is why Prasad called for interdisciplinary teams of engineers, lawyers, ethicists, and social scientists.

—

3.2 Synthesized Content (Claude API Output)

The Claude API synthesized the Perplexity research into a 1800-word article structured as follows:

3.2.1 Section A: Why Privacy Matters in AI

Privacy is not an afterthought in artificial intelligence. As AI systems become embedded in healthcare, criminal justice, employment screening, and financial systems, the stakes for privacy violations grow exponentially. When AI misuses personal data, the harm is direct: individuals face discrimination, surveillance, identity theft, and loss

of autonomy. Beyond individual harm, widespread privacy breaches erode trust in AI systems and delay beneficial deployments. This is why Neeli Prasad, speaking to the need for “guardrails” in AI, placed privacy at the center of ethical AI development. It must be built in from day one.

3.2.2 Section B: Real Harms and the Cost of Privacy Failure

The risks are not theoretical. In 2023, a Hong Kong bank lost millions when a deepfake video and voice, impersonating the CEO, authorized a fraudulent transfer. Healthcare providers and insurance companies have had patient records breached, exposing sensitive data that was then used for blackmail or sold. Children have been harassed with AI-generated fake images and voice recordings, with some cases linked to documented suicides. These aren’t edge cases. Without privacy protections built in, AI causes harm.

3.2.3 Section C: Regulatory Frameworks

The EU AI Act is the first law written specifically for AI systems. It mandates privacy impact assessments for high-risk applications like facial recognition, employment screening, and criminal justice algorithms. It requires transparency, explainability, and human oversight. The GDPR, passed ten years earlier, covers baseline data protection: data minimization, purpose limitation, explicit consent, and the right to explanation. But GDPR wasn’t designed for how AI actually works. Modern AI combines massive datasets and repurposes them in ways the original data subjects never imagined.

3.2.4 Section D: Technical Safeguards

Privacy-by-design principles call for privacy to be engineered into systems from the architecture stage, not bolted on afterward. Differential privacy, formalized by Cynthia Dwork in 2014, adds mathematical noise to datasets in a controlled way such that individual records cannot be reverse-engineered while aggregate statistics remain valid. Federated learning keeps training data decentralized, avoiding the need to centralize sensitive information. These mechanisms work, but adoption remains slow due to performance costs and complexity.

3.2.5 Section E: Governance and the Role of Human Oversight

Technical solutions alone are insufficient. Privacy protection requires lawyers, ethicists, social scientists, and engineers working together from the start of design. Prasad

called for this interdisciplinary collaboration. Humans can't be passive observers. They must understand how AI systems use data. They need the ability to override AI decisions. And they must exercise real control over personal information. Transparency in AI tool use, what Prasad called "prompt engineering" as a discipline, extends to governance too. AI systems need to explain their data practices to the people affected by them.

3.3 Writing Style Comparison (QuillBot Output)

To explore how writing style affects clarity, the following paragraph was processed through QuillBot's three modes:

3.3.1 Original (Claude API)

"Privacy is a fundamental right that must be protected in artificial intelligence systems. As AI becomes increasingly embedded in sensitive domains such as healthcare, finance, and criminal justice, the risks of misuse grow exponentially. Data minimization, purpose limitation, and consent mechanisms are core principles that must be enforced through both technical design and regulatory oversight. Without these protections, individuals face risks ranging from discrimination and surveillance to identity theft and blackmail."

3.3.2 Standard Mode (QuillBot)

"Safeguarding privacy is essential in AI systems. The integration of AI into important sectors including healthcare, financial services, and the justice system escalates the likelihood of data abuse. Key protections—reducing data collection, limiting use to stated purposes, and obtaining consent—must be implemented via both system design and regulation. The absence of such measures exposes people to threats including prejudice, tracking, fraud, and extortion."

3.3.3 Formal Mode (QuillBot)

"The protection of privacy constitutes a fundamental imperative within artificial intelligence systems. Given the expanding deployment of artificial intelligence technologies in sensitive domains—including healthcare provision, financial services, and

criminal justice administration—the potential for data misuse is amplified considerably. Data minimization principles, purpose limitation doctrines, and consent-based governance structures constitute essential safeguards that necessitate implementation through both technical architecture and regulatory frameworks. Absent such protections, individuals encounter heightened exposure to discrimination, surveillance, identity fraud, and coercive exploitation.”

3.3.4 Analysis

- **Original:** Balanced complexity; suitable for academic but not overly formal; accessible to the target audience (engineering students).
- **Standard:** More concise (70 words vs. 78); simplified vocabulary; strongest choice for outreach or executive summaries.
- **Formal:** Increased complexity and latinate vocabulary; 85 words; likely appropriate for peer-reviewed publication but potentially reduces accessibility for undergraduates.

All three convey the same core message. We kept the Original (Claude API) version because it’s precise and clear enough for engineering students to understand.

3.4 Presentation Output (SlidesAI.io)

SlidesAI.io generated a 9-slide presentation deck with the following structure:

1. **Title Slide:** “Privacy and Data Protection in AI Systems”
2. **Agenda:** Overview of the 5 main sections
3. **Why Privacy Matters:** Key statistics on data breaches, cost of privacy failures
4. **Real-World Cases:** Hong Kong deepfake fraud, patient data breaches, cyberbullying
5. **Regulatory Landscape:** GDPR and EU AI Act requirements and timelines
6. **Technical Safeguards:** Privacy-by-design, differential privacy, federated learning

7. **Governance and Transparency:** Interdisciplinary collaboration, human oversight
8. **Key Takeaways:** Privacy is foundational, requires both technical and governance solutions
9. **Conclusion Slide:** Call to action for responsible AI development

The presentation was generated in approximately 12 minutes and required no manual editing. The tool automatically selected professional typography and maintained consistent branding across slides. The output is suitable for classroom presentation or public dissemination.

3.5 Comparative Analysis: Tool Quality and Efficiency

Tool	Output Quality	Execution Time	Cost	Suitability
Perplexity	4/5	12 min	Free	Research foundation
Claude API	5/5	4 min	Free	Primary synthesis
QuillBot	3/5	9 min	Free (limited)	Style exploration
SlidesAI.io	3.5/5	12 min	Free (basic)	Rapid presentation

Table 2: Comparative quality and efficiency assessment. Quality rated on accuracy, completeness, and relevance (1–5 scale).

3.5.1 Quality Assessment

- **Perplexity:** High quality in source curation and citation; minor limitation in search depth without premium subscription.
- **Claude API:** Highest overall quality; synthesized content was coherent, well-structured, and demonstrated deep understanding of both technical and governance dimensions of privacy.
- **QuillBot:** Effective for style exploration but limited by free-tier word count (125 words per paraphrase); required multiple passes for longer text.
- **SlidesAI.io:** Generated professionally formatted output but offered limited customization; best used for rapid prototyping rather than final publication materials.

3.5.2 Efficiency

The complete pipeline (research $\rightarrow$ synthesis $\rightarrow$ refinement $\rightarrow$ presentation) was executed in approximately 45 minutes of active time. The combination of free and freemium tools made the process accessible to students without budget constraints. Total cost: \$0.

3.5.3 Observed Bottlenecks

- QuillBot’s 125-word limit required splitting longer paragraphs into multiple paraphrase operations.
 - SlidesAI.io’s template library limited design flexibility for specialized content (e.g., mathematical formulas, detailed diagrams).
 - Perplexity’s free tier returned 3–5 sources per query; premium versions would provide more comprehensive coverage.
-

3.6 Pipeline Effectiveness

Running multiple tools in sequence worked. Perplexity gathered credible research. Claude synthesized it into a coherent narrative. QuillBot showed how changing writing style affected clarity. SlidesAI.io generated slides quickly. No single tool did everything. The value came from combining them. Prasad made this point: future AI practitioners need to be problem solvers who can combine multiple tools and disciplines.

4 Conclusions

The pipeline worked. We ran four free tools in sequence, 45 minutes of work, and produced a complete knowledge artifact on Privacy and Data Protection in AI Systems. This is a topic Prasad said matters for responsible AI development.

4.1 Key Accomplishments

The pipeline delivered a complete, multi-format output in 45 minutes using only free and freemium tools. Claude was the strongest link, producing clear, well-structured

writing that engineering students could read. Each tool brought something different to the work: Perplexity’s citations, Claude’s analysis, QuillBot’s style options, SlidesAI.io’s design. That variety mattered.

4.2 Limitations and Trade-Offs

The free tier of each tool imposed practical constraints. QuillBot’s 125-word paraphrase limit required multiple passes for longer text. SlidesAI.io’s template library, while professional, offered limited flexibility for specialized content. Perplexity’s search depth, without premium subscription, returned a smaller citation pool than would be ideal for a comprehensive literature review. These constraints are not failings of the tools themselves but rather reflect the reality that production-quality outputs often require either paid subscriptions or manual refinement.

4.3 Broader Lessons on Human-AI Collaboration

This project illustrates three lessons that directly connect to Prasad’s seminar themes:

1. **Interdisciplinary Orchestration:** Just as Prasad called for engineers, lawyers, ethicists, and social scientists at the “round table” during AI governance, the tool pipeline showed that different AI tools serve different roles. No single tool was sufficient; value emerged from composition.
2. **Transparency and Prompt Engineering:** By documenting every prompt, configuration, and output, this project followed Prasad’s insistence that students “disclose your prompt engineering.” Documentation enables reproducibility. It also trains practitioners to think critically about the inputs and parameters that shape AI outputs. Show the work, not just the result.
3. **Human Judgment and Oversight:** While the tools accelerated the creation process, human decisions were essential at every stage: evaluating research quality, assessing synthesis accuracy, choosing which style variant was most appropriate, and determining when an output was ready for dissemination. The tools amplified human capability; they did not replace it.

4.4 Implications for Privacy in AI

Privacy and data protection in AI are urgent. The EU AI Act is now in effect. Students and engineers entering the field in 2026 will work in environments where privacy is not optional. This project practiced skills that matter: pulling together

regulatory requirements and technical safeguards, then explaining them to different audiences. Modern AI tools let practitioners move beyond purely technical work into broader roles that integrate law, governance, and social impact.

4.5 Future Work

Possible extensions of this project include:

- Integration of additional tools (e.g., Gemini for longer synthesis, Canva for custom presentation design) to compare quality across different providers.
- A longitudinal comparison: rerunning this pipeline in 12 months to assess how evolving AI capabilities and regulatory changes affect the outputs.
- Peer review: having the synthesized content reviewed by privacy experts to assess not just stylistic quality but factual accuracy and completeness.
- Application to other seminar topics: demonstrating that this methodology is generalizable across different subject domains.

4.6 Final Reflection

Prasad closed with "lead with integrity, not just innovation." We tried to live that in three concrete ways. We were transparent about using AI tools, showing the prompts and outputs. We grounded the work in real privacy failures—not hypothetical harms. And we made human judgment central at every step, not as an afterthought.

The tools were powerful and fast. But a tool is just a tool. The hard part is knowing what to ask for and when to say no to what it offers. That judgment is where humans still matter.

5 References

- [1] European Parliament and Council. Regulation (eu) 2016/679 — general data protection regulation. <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:32016R0679>, 2016. Official Journal of the European Union, L 119.
- [2] European Parliament and Council. Regulation (eu) 2024/1689 — artificial intelligence act. <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:32024R1689>, 2024. Official Journal of the European Union, L 188.
- [3] Neeli Prasad. Human-ai collaboration: Ethics, privacy, and governance. Universidade Estadual de Campinas, School of Electrical and Computer Engineering, May 2026. IEEE Distinguished Lecture, IA382A Seminar Series.